

ATRIO

Your AI boardroom.

Six specialist agents. One audited table. Mandate-enforced treasury.

The problem

Founders and family offices decide alone.

Big decisions go one of two ways today:

- Delegated to a single advisor — fast, but a single point of failure.
- Convened with a committee — slow, hard to schedule, hard to audit.

No middle option. No audit trail of the reasoning. No way to enforce a mandate at machine speed. No way to replay a decision in six months.

The solution

An AI boardroom that holds a real debate.

DEBATE

Six specialist agents (CFO, CTO, CMO, COO, Counsel, Facilitator) argue with each other before answering. Dissent triggers re-runs.

ENFORCE

Per-tenant mandate is checked at the API on every treasury action. Two-party authorisation cannot be bypassed by any client.

AUDIT

Append-only audit log captures every turn, vote, model invocation, and state transition. Exportable as JSONL + manifest for compliance.

Six specialists

Distinct personas. Distinct models. One table.

CFO Capital + burn <i>Gemini 2.5 Flash</i>	CTO Tech feasibility <i>Featherless Qwen2-72B</i>	CMO Narrative + GTM <i>Featherless Llama-3.1-70B</i>
COO Ops + people <i>Featherless Mistral-Large</i>	Counsel Legal + regulatory <i>Gemini 2.5 Pro</i>	Facilitator Synthesis + dissent <i>Gemini 2.5 Pro</i>

Model registry is the only path to inference. Fallback chains are configured per agent (Gemini primary, Featherless fallback, mock for offline demo). Every model invocation writes to the audit log with token counts, latency, and request ID.

Treasury · Mandate · Two-party

The agents can do, not just advise.



Four mandate gates checked at the API layer (not the UI):

✓ Permitted instrument list

SHV-xStock, IEF-xStock, EURUSD-xStock

✓ Permitted side

buy / sell

✓ Single-instrument max

€50,000 per proposal

✓ Daily loss limit

€25,000 rolling 24h

✗ Two-party authorisation cannot be bypassed.
The proposer cannot self-second. Tested — and demonstrated live in the recording.

Live demo · 4 stages

What you will see in the 2-minute video.

1

BOARDROOM DEBATE

Founder asks a Q3 hiring question; PDF burn-plan uploaded; six agents debate; dissent re-runs; consensus + action list.

2

TREASURY · MANDATE · TWO-PARTY

Founder proposes SHV-xStock buy. Founder authorises (1/2). Founder TRIES to self-second — REFUSED by the API.

3

SECOND HUMAN

CEO signs in on a separate browser, same tenant. Authorises (2/2). Trade executes against Kraken paper. Audit gains 5 rows.

4

BOARDPACK · AUDIT EXPORT

Session closed. Boardpack PDF download. Audit page. Export ZIP = JSONL + manifest. Compliance-ready.

Proof, not promises

Audit-grade comes from tests, not slogans.

381 / 381

backend tests pass

90.68 %

line coverage (gate 85 %)

15 / 15

vitest pass · typecheck clean

16 / 20

Playwright pass against live stack

24 / 24

verification-a · hard API assertions

14 / 14

verification-b · OCR rubric items

12

real bugs found and fixed

23

commits on main · all green

ALL numbers reproducible on the repo. ``make up && pytest && pwsh ./demovideo/verification-a/run-verify-a.ps1``

Architecture

Boring, predictable, audit-grade.

FRONTEND

React · Vite · Tailwind

Caddy SPA on :8080 (PWA-ready)

API

FastAPI · SQLAlchemy · Alembic

uvicorn on :8000 (RLS · mandate gates)

INFRA

Postgres 16 + pgvector · MinIO · LiveKit · Mailhog

All in compose · one ``make up``

Sponsor pools wired: Vultr (deploy) · Gemini (inference) · Featherless (fallback) · Speechmatics (voice) · Kraken (treasury paper)

Sponsor pools utilised

Five pools. Real integration. No demoware.

Vultr	Compute + deployment	<i>VM target for the public demo URL</i>
Google Gemini	Primary inference (CFO + Counsel + Facilitator)	<i>API key configured · model registry routes calls</i>
Featherless	Specialist fallback (CTO + CMO + COO)	<i>Configured · activates when Gemini is down or rate-limited</i>
Speechmatics	Live STT + diarisation (voice mode)	<i>Sidecar wired · text-only fallback in demo recording</i>
Kraken	Treasury execution (xStocks)	<i>Paper-mode by default · live mode is one config flip</i>

Beyond the MVP

What ships next.

v1.1 — 3 weeks

- WCAG 2.1 AA via [@axe-core/playwright](#)
- Performance gates (p95 < 10s session)
- Speechmatics live STT (coupon redeemed)
- Featherless live fallback chain
- Push backend coverage 90.68 % → 100 %

v1.2 — 8 weeks

- Native mobile app (PWA → Capacitor)
- IOTA anchoring on treasury receipts
- Vultr Object Storage (from MinIO)
- Multi-language UI (EN / IT / ES / FR / DE)
- Long-horizon memory (per-agent profile)

Team + ask

Built by Verixa for Milan AI Week 2026.

Founder + lead engineer · architecture, infra, demo, deployment

Generated under the AT-Hack0021 sprint with Claude as paired engineer.

What we are asking for

- Vultr credits to host the live demo URL for the EU founder cohort.
- Gemini + Featherless production-tier keys for the v1.1 multilingual rollout.
- 30 minutes with a Kraken xStocks PM to validate the mandate-violation corpus.
- Time on stage at Milan AI Week to demo end-to-end with audience-supplied questions.

ATRIO Boardroom

Mandate-enforced at the API.

Two-party authorisation. Cross-tenant isolation. Audit-grade by default.

github.com/vsenthil7/atrio-boardroom

Apache 2.0 · 381 backend tests · 90.68 % coverage · 24/24 + 14/14 demo-verified

Thank you.